PHASE 1 PILOT — BENEFITS CASE

Reusing MITR AI to Deliver the Phase 1 Pilot

An already-built conversational analytics UX for Revenue, Cost, General Ledger, and Consolidation

Pilot value in weeks, not quarters — the same surface extends into Phase 2 and Phase 3

Companion to the Architecture Review Board deck · Finance Engineering · September 2026

WHY THIS, NOW

The Phase 1 Pilot needs a vehicle — one already exists

The roadmap calls for an AI-enabled UX plus 1–2 agents on the highest-pain stream, with a human-in-the-loop review workflow and a controls checkpoint.

- **Built, not hypothetical** — MITR AI already does plain-language querying, generated insights, and governed data access across finance data. It is demoable today.
- **Wrap, don't replace** — it reads through governed feeds and sits above the ledgers. No change to systems of record or controls — the ARB's own approach.
- **Weeks to first user** — configuration plus one data connector, against the two to three quarters a pilot UX built from scratch would take.
- **Not throwaway** — the same chat and dashboard surface is where Agent One's drafting agents plug in for Phase 2.

Bottom line — *piloting on MITR AI tests the UX and the controls model on real finance data before the larger build commits.*

WHAT IT LOOKS LIKE

Seeing it: the pilot surface today

Live screens from the running platform — detail on the slides that follow.

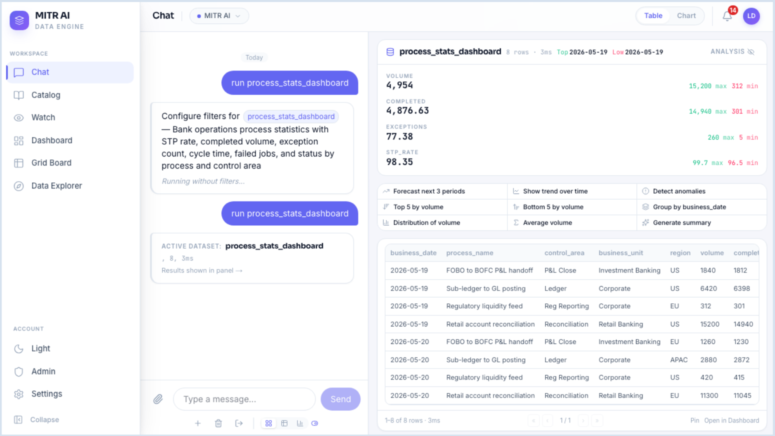

Ask in plain language

Data, KPI tiles with deltas, and one-click analysis — no SQL, no report to build.

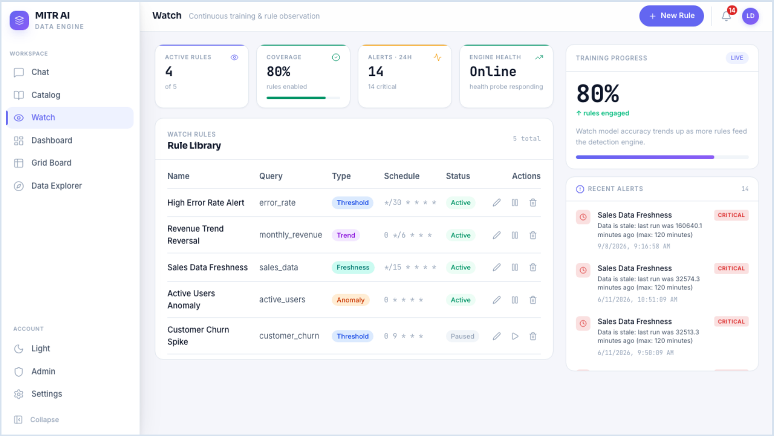

Monitors that run themselves

Freshness, threshold, trend, and anomaly rules raise alerts before anyone asks.

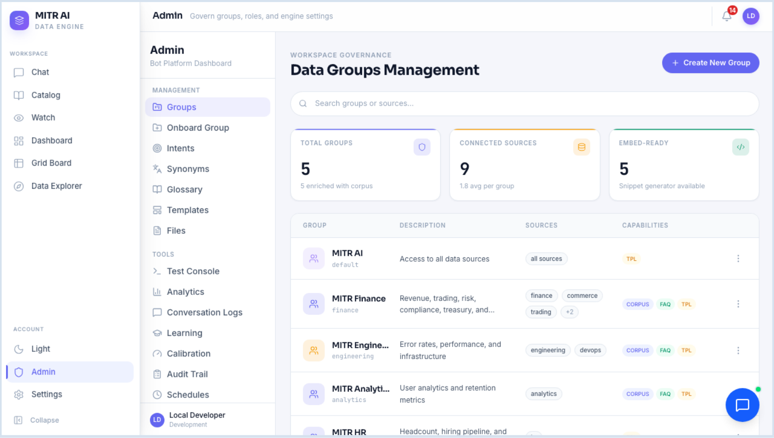

Governed from one console

Per-group data scoping, connectors, vocabulary, learning, and a full audit trail.

Captured from the running MITR AI build, using a mock finance dataset.

THE SURFACE

A question in, data out — no SQL, no report build

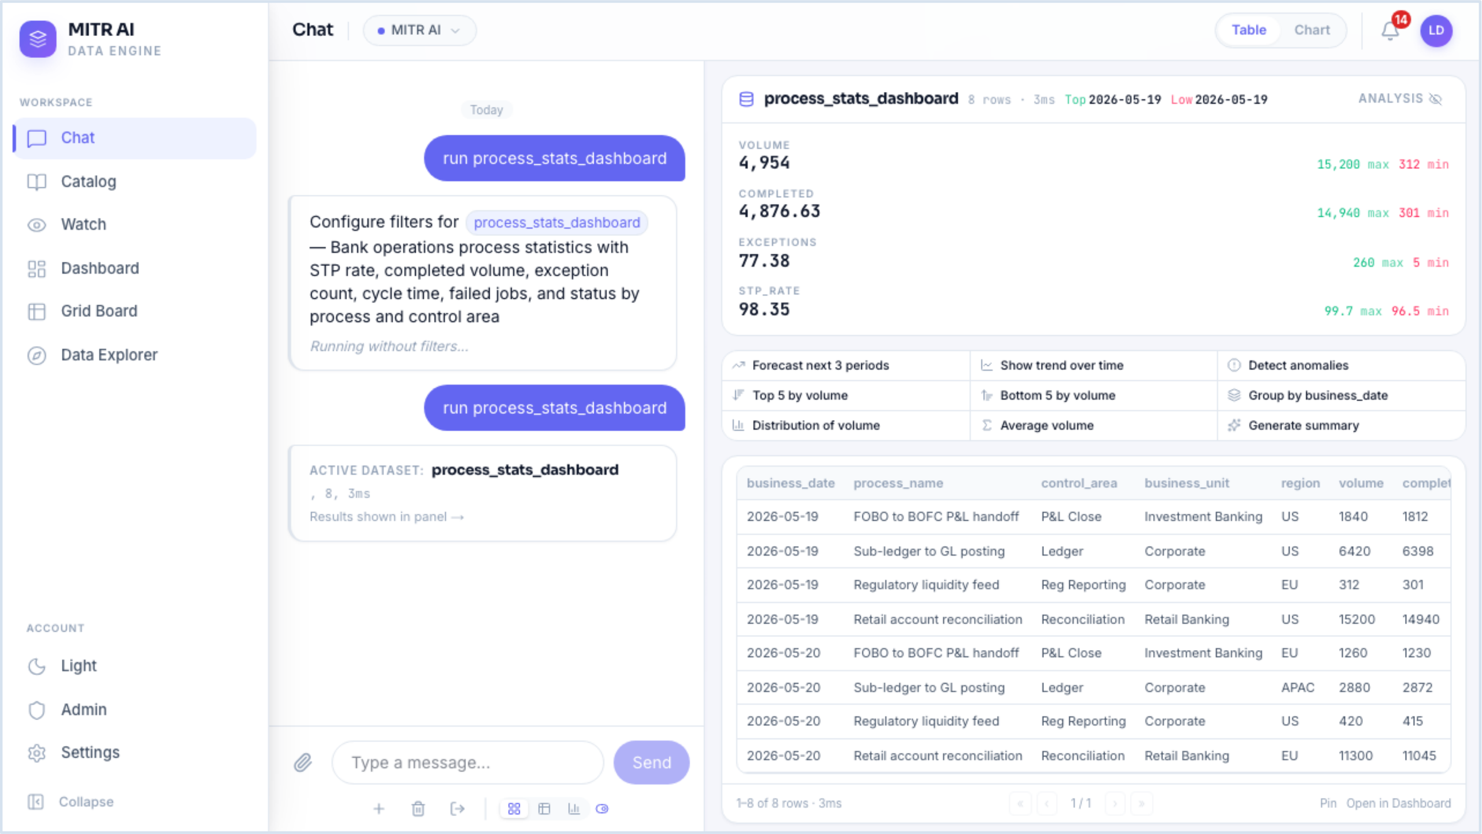

Chat → run process_stats_dashboard, filters skipped.

WHAT TO NOTICE

- Plain-language query returns a result set in about three seconds.
- KPI tiles are computed automatically, each with its max / min range.
- One click runs a forecast, trend, anomaly scan, group-by, or summary.
- Table, chart, and drill-down live in the same view — nothing to export.

ANALYTICS LAYER

Insight appears before anyone asks for it

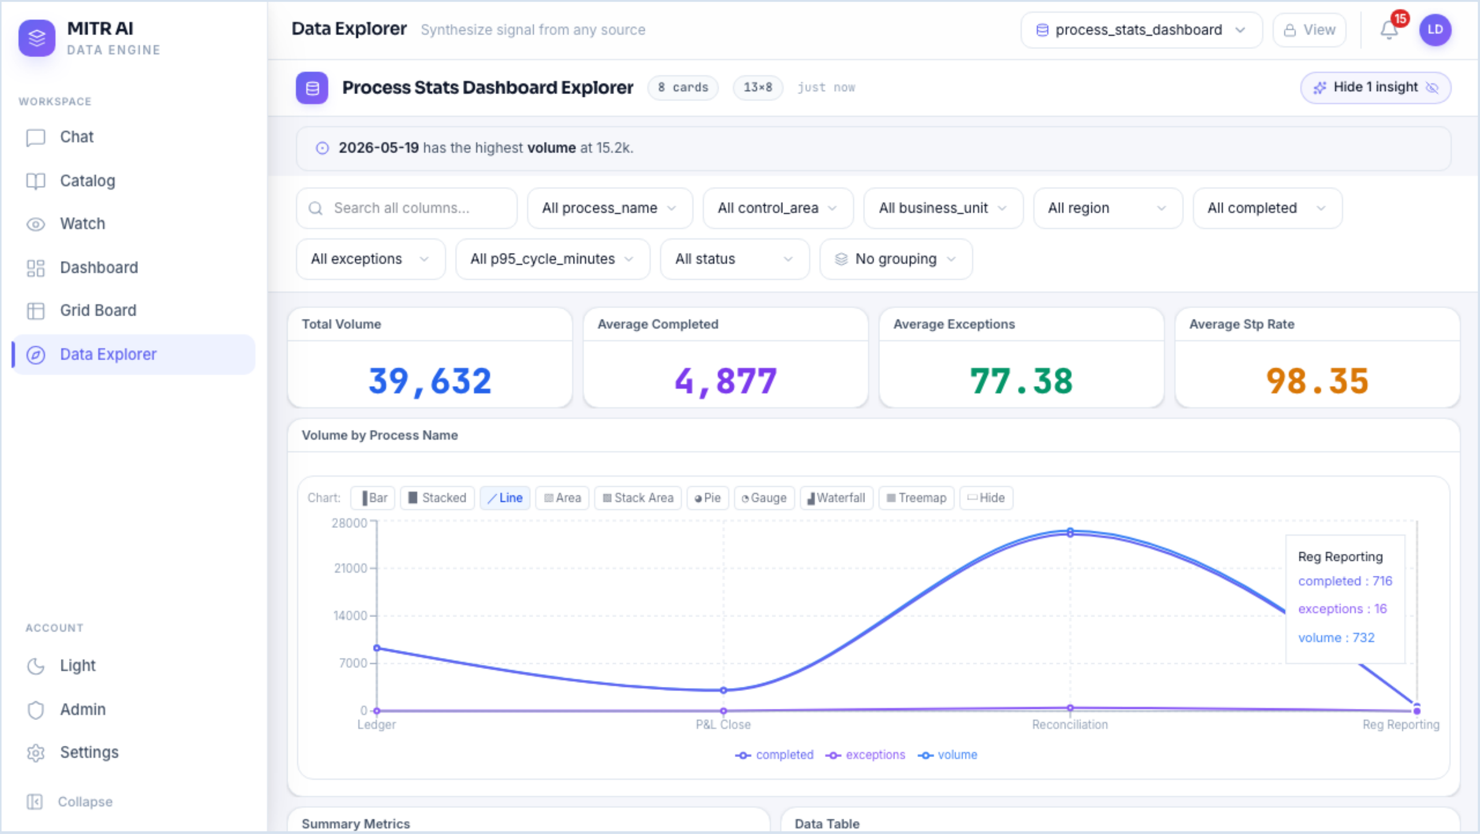

Data Explorer → process_stats_dashboard.

WHAT TO NOTICE

- The top insight is surfaced on load — “2026-05-19 has the highest volume at 15.2k.”
- KPI cards show direction and % change against the prior period.
- Interactive chart, eight chart types, click any point to drill in.
- Deterministic — no LLM in the path, every number is traceable to source.

The same before / after — most of it working today

The ARB's target experience, mapped to what the pilot delivers now versus what Agent One adds later.

TODAY	WITH MITR AI (PILOT)	STATUS
A user pulls a variance report and scans for outliers by hand	Flagged variances are surfaced automatically, each with a plain-language explanation of the likely driver	Live today
"What's driving the movement?" takes an analyst hours	Root-cause ranking by share of total, returned in the chat in seconds	Live today
Cross-stream reporting means exporting and combining data by hand	One semantic layer — cross-stream metrics are already assembled and ready	Live today
Stale or broken data is noticed only when someone stumbles on it	Anomaly and freshness monitors raise it proactively, with a full audit trail	Live today
Journal entries and eliminations are typed from source documents	Agents draft the entry from source data; a preparer reviews and approves	Phase 2 same surface

The platform flags the outlier and explains it

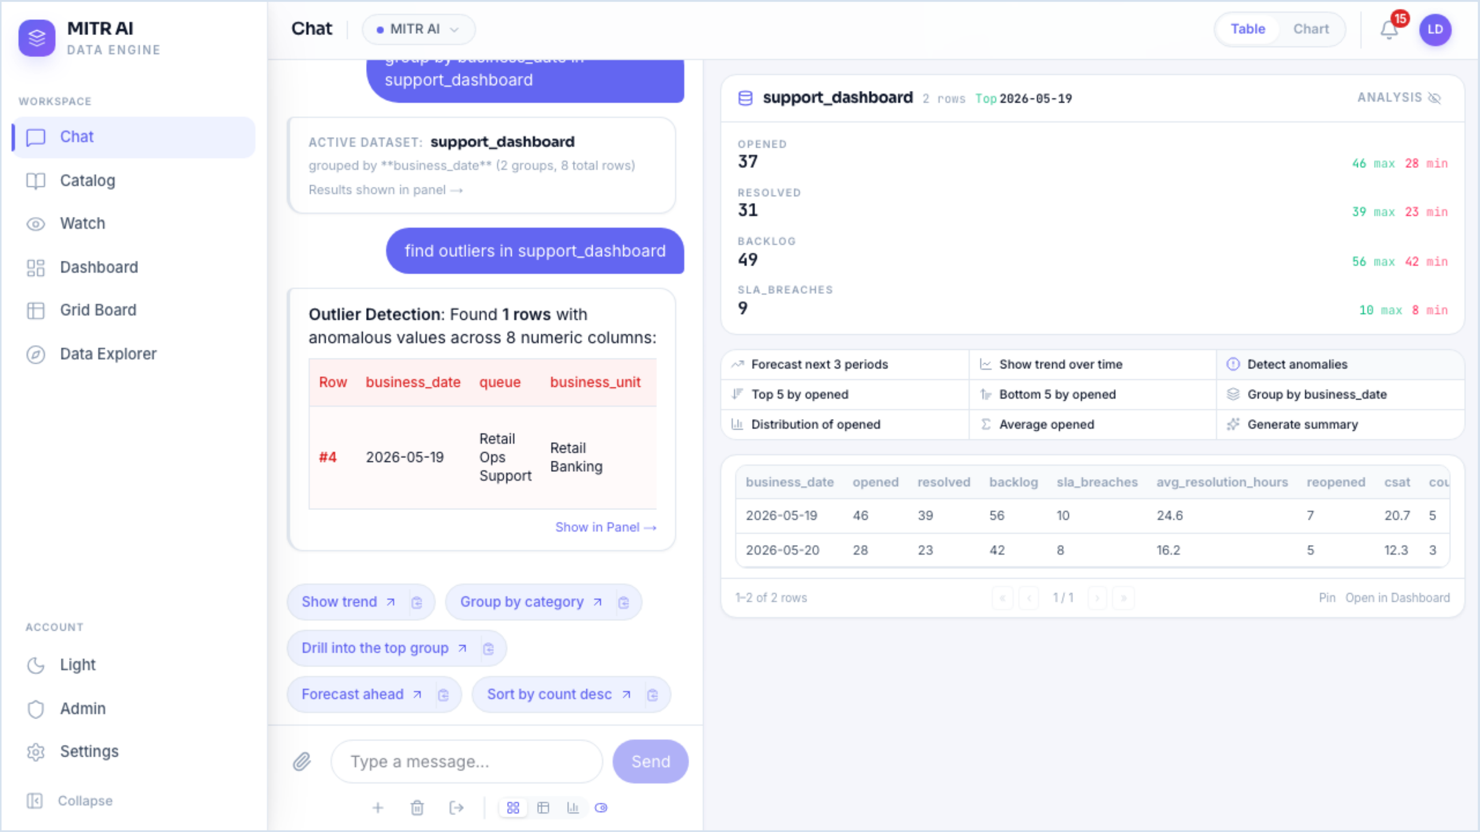

Chat → "find outliers in support_dashboard."

WHAT TO NOTICE

- Outlier detection runs across every numeric column, unprompted.
- Plain-language reason: "Resolved 16 vs avg 7.75 — unusually above average."
- The same mechanism drives variance explanation across all four streams.
- A preparer reviews the finding; nothing is hidden or auto-posted.

HOW IT MAPS

MITR AI already delivers the UX and Analytics layers

Same layer model. The pilot delivers the UX layer and the Analytics layer — a named differentiator — and proves the review workflow Agent One will use.

UX Layer	**PROVIDES** — chat, live dashboards, and grid board; AI insight and automation status surfaced by default
Agent Orchestration — Agent One	**PHASE 2** — out of pilot scope; the pilot proves the human-in-the-loop review workflow agents post through
Analytics Layer	**PROVIDES** — variance narratives, anomaly detection, and root-cause — deterministic, no LLM, fully auditable
EDP / Data Products (read-only)	**CONSUMED** — SQL Server, Oracle, CSV, and REST connectors — read-only, health-monitored, circuit-broken
Integration / Data Layer	**PROVIDES** — per-group data scoping — the only path the platform uses to reach source data
Systems of Record	**UNCHANGED** — Revenue, Cost, GL, and Consolidation sub-ledgers — untouched by the pilot

Guardrails already in the product — full audit trail for every action · role-based access across the UX layer · PII detection and masking with per-view audit · read-only data access enforced at the connector · end-to-end request correlation IDs. Aligns to SOX and existing control frameworks.

CONTROLS

Freshness and anomaly rules run on their own

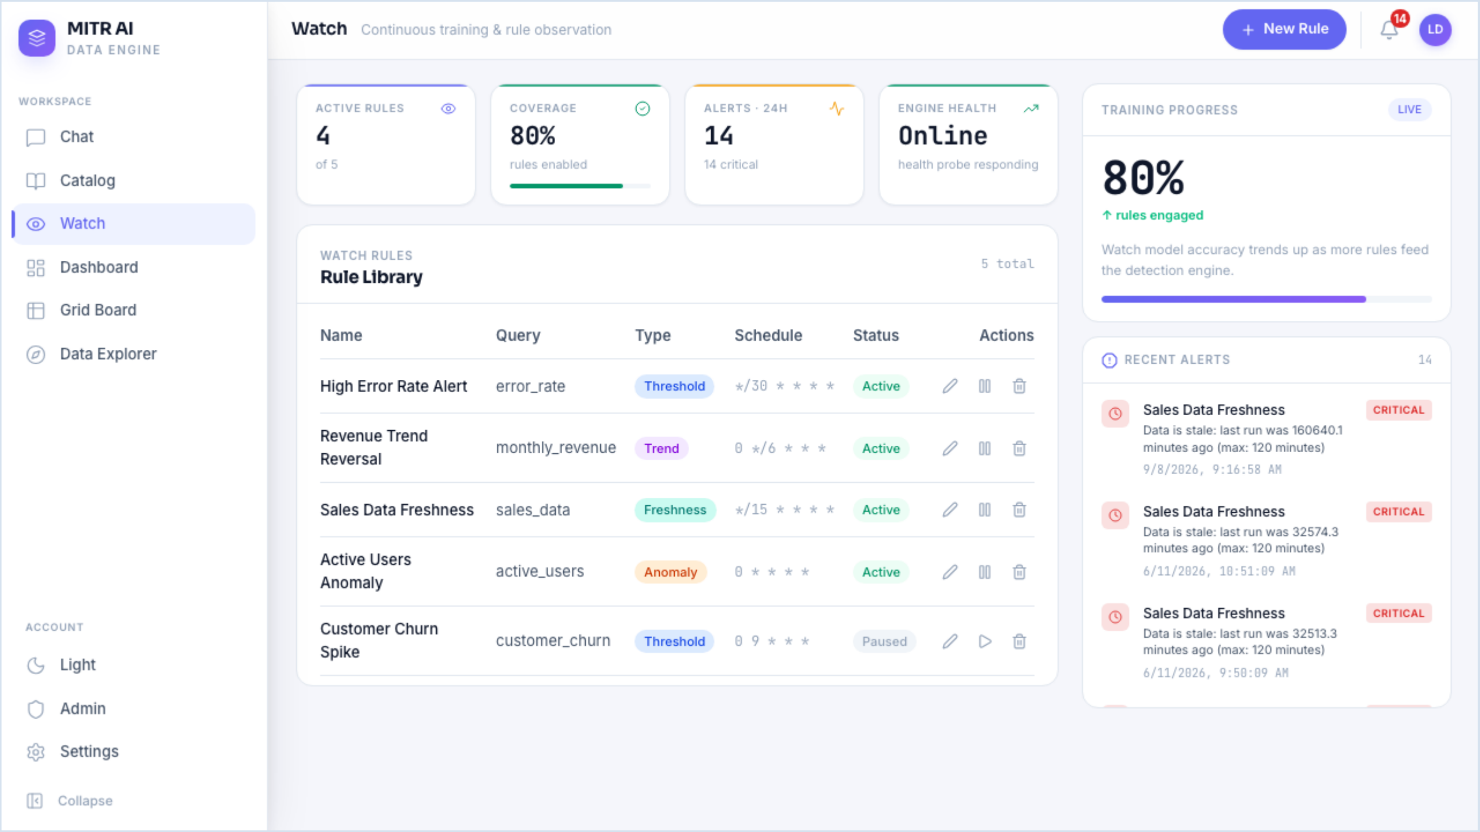

Watch → rule library and recent alerts.

WHAT TO NOTICE

- Rule types cover threshold, trend, freshness, and statistical anomaly.
- Checks are scheduled — alerts are raised with no one watching the screen.
- Every alert is timestamped and severity-tagged for the audit trail.
- Rule coverage and detection-model training are tracked over time.

GOVERNANCE

Data scoping, connectors, and vocabulary in one place

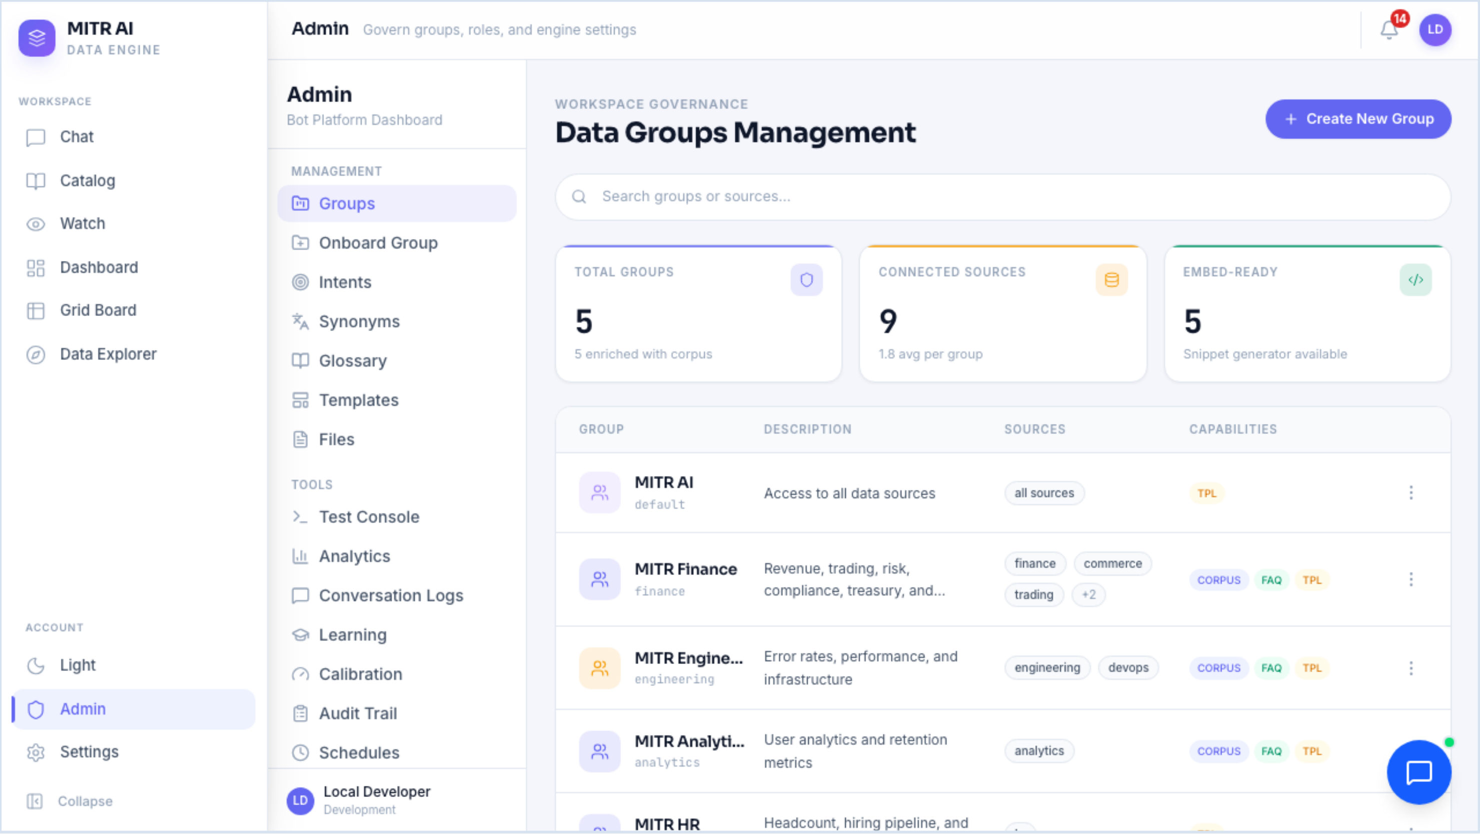

Admin → Data Groups Management.

WHAT TO NOTICE

- Per-group data access — five groups over nine connected sources.
- Capabilities are set per group (corpus, FAQ, templates).
- Intents, synonyms, glossary, learning queue, and calibration all live here.
- Full audit trail and conversation logs for every group.

DECISION NEEDED

What I’m asking

	Build the pilot new	Pilot on MITR AI
Time to first finance user	Two to three quarters	**Weeks**
Cost	A new pilot build team	**Config + one connector on an existing codebase**
Risk	Unproven UX and controls	**150+ automated checks, load-tested, demoable now**
Longevity	Pilot code may be thrown away	**Same surface carries into Phase 2–3**

THE ASK

- Endorse MITR AI as the Phase 1 Pilot vehicle for the AI-enabled analytics UX.
- Help confirm the pilot stream and a business sponsor to co-own the success criteria.
- Approve a short spike to connect it to one real finance data product.

Success measured by: close-cycle time reduction • manual rework reduction • % of insights and entries AI-originated • adoption across finance teams.